Data Science & AI Bootcamp – Interview Question Bank

GENERATIVE AI

1. How does Lang Chain help with Gen AI Application development?
2. Explain different architectures (GPT, BERT, T5).
3. How do you address the context window length limitation while developing Gen AI Application?
4. What is the role of a vector database in a Generative AI pipeline?
5. How do transformers work?
6. Explain RLHF (Reinforcement Learning from Human Feedback).
7. Why do we need Generative AI and why can't we handle the use case using Deep Learning?
8. How do you implement model distillation?
9. What are LLMs used for?
10. What are language models?
11. How do you handle model biases?
12. How and why are transformers better than RNN architectures?
13. How do you implement content filtering?
14. How do you implement model quantization?
15. Explain tokens and tokenization.
16. Explain text generation techniques.
17. What is semantic search?
18. Explain vector databases.
19. What is retrieval-augmented generation?
20. How do you implement fine-tuning for LLMs?
21. What are parameter-efficient fine-tuning methods?
22. What is fine-tuning in LLMs?
23. What are the different types of fine-tuning?
24. What is temperature in LLMs?
25. When to use Prompt Engineering, RAG, or Fine-Tuning?
26. What is few-shot learning?
27. Explain prompt templates.
28. How do you handle model drift?
29. Explain prompt chaining strategies.
30. What are instruction-tuned models?
31. What are different evaluation frameworks?
32. What are different model compression techniques?
33. What is prompt engineering?
34. What are different types of LLMs?

Data Science & AI Bootcamp – Interview Question Bank

35. Explain different deployment strategies.
36. How do you evaluate generative models?
37. What are different prompting strategies?
38. How do you handle model hallucinations?